

Understanding Heart Disease:

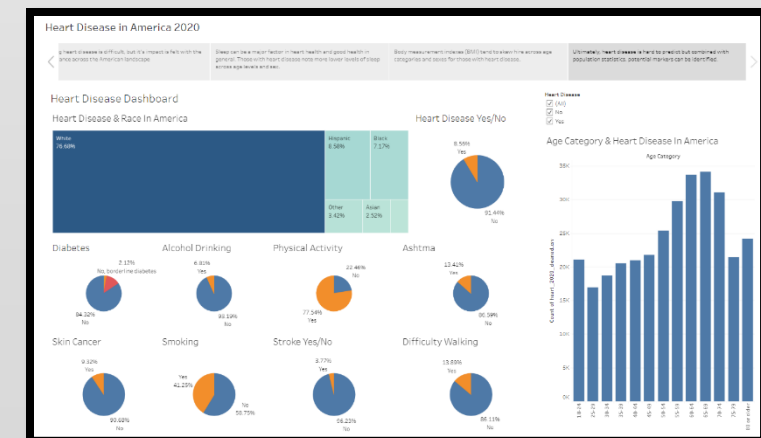
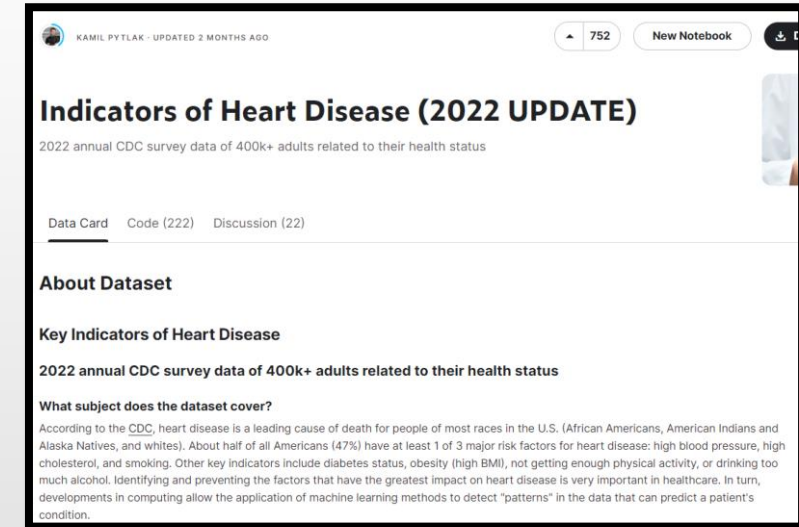
Can You Predict Who Gets It?

OLIVER KISZA, SPENCER AUSLANDER,
VIGNESH CHERIATH, BRANDON MATA,
FEMI ADEMUWAGUN, CHRIS MANFREDI



Data Retrieval, Multivariate Regression

- Data pulled from Kaggle data sets (data originates from CDC)
- Dataset originally comes from the CDC and is a major part of the Behavioral Risk Factor Surveillance System (BRFSS)
- Initially looked at both data sets (2020, 2022) but went with 2020 due to data simplicity and usage of Heart Disease (Yes/No) as dependent variable
- Initial Data Cleaning was required to use SPSS as our initial data variable understanding
- Regressions were not very powerful but did yield ideas around important variables to use and look through in the population



Data Retrieval, Multivariate Regression

2020 – Heart Disease

Model	Variables Entered	Variables Removed ^a
1	GeneralHealth	
2	Stroke	
3	Diabetic	
4	KidneyDisease	
5	DifficultyWalking	

Model Summary ^p				Sex
Model	R	R Square	Adjusted R Square	
1	.254 ^a	0.065	0.065	Cancer
2	.307 ^b	0.094	0.094	Smoking
3	.328 ^c	0.107	0.107	Days
4	.340 ^d	0.115	0.115	Days
5	.348 ^e	0.121	0.121	BMI
6	.357 ^f	0.127	0.127	Smoking
7	.363 ^g	0.132	0.132	Sleep
8	.367 ^h	0.134	0.134	Asthma
9	.368 ⁱ	0.135	0.135	Activities
10	.369 ^j	0.136	0.136	
11	.370 ^k	0.137	0.137	
12	.370 ^l	0.137	0.137	
13	.370 ^m	0.137	0.137	
14	.371 ⁿ	0.137	0.137	
15	.371 ^o	0.137	0.137	

2022 – High Risk

Model	Variables Entered	Variables Removed ^a
1	HIVTesting	
2	DifficultyConcentrating	
3	AlcoholDrinkers	
4	HadArthritis	
5	HadDepressiveDisorde	
6	FluVaxLast12	
7	Sex	

Model Summary ^p				Walking
Model	R	R Square	Adjusted R Square	
1	.130 ^a	0.017	0.017	Cancer
2	.151 ^b	0.023	0.023	
3	.171 ^c	0.029	0.029	
4	.182 ^d	0.033	0.033	
5	.191 ^e	0.036	0.036	
6	.197 ^f	0.039	0.039	
7	.201 ^g	0.040	0.040	Hands
8	.204 ^h	0.041	0.041	OfHearing
9	.205 ⁱ	0.042	0.042	Health
10	.206 ^j	0.042	0.042	
11	.207 ^k	0.043	0.043	
12	.207 ^l	0.043	0.043	
13	.208 ^m	0.043	0.043	Ever
14	.208 ⁿ	0.043	0.043	Activities
15	.209 ^o	0.044	0.044	PressingBathi
16	.209 ^p	0.044	0.044	Attack
17	.209 ^q	0.044	0.044	onDifficulty
18	.210 ^r	0.044	0.044	
19	.210 ^s	0.044	0.044	
20	.210 ^t	0.044	0.044	
21	.210 ^u	0.044	0.044	
22	.210 ^v	0.044	0.044	
23	.210 ^w	0.044	0.044	

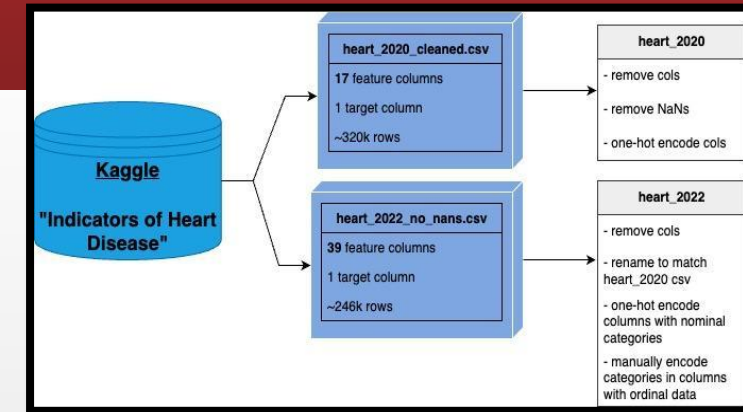
2022 – Heart Attack

Model	Variables Entered	Variables Removed ^a
1	HadAngina	
2	HadStroke	
3	GeneralHealth	
4	ChestScan	
5	Sex	
6	DifficultyWalking	

Model Summary ^p				PD
Model	R	R Square	Adjusted R Square	
1	.455 ^a	0.207	0.207	Drinkers
2	.470 ^b	0.221	0.221	VaxEver
3	.477 ^c	0.227	0.227	HardOfHearing
4	.481 ^d	0.231	0.231	neyDisease
5	.483 ^e	0.233	0.233	
6	.485 ^f	0.235	0.235	Activities
7	.486 ^g	0.236	0.236	VisionDifficulty
8	.486 ^h	0.236	0.236	hma
9	.487 ⁱ	0.237	0.237	ritis
10	.487 ^j	0.237	0.237	ressiveDisorder
11	.488 ^k	0.238	0.238	kLastYear
12	.488 ^l	0.238	0.238	ck
13	.488 ^m	0.238	0.238	
14	.488 ⁿ	0.238	0.238	
15	.488 ^o	0.238	0.238	
16	.488 ^p	0.239	0.239	
17	.488 ^q	0.239	0.239	
18	.488 ^r	0.239	0.239	

Python Data Cleaning

2 Datasets: 2020 And 2022(updated) Annual CDC Survey Data Of 400,000+ Adults



original file (heart_2020_cleaned.csv):

- 17 feature columns (heart disease indicators)
- 1 target column (HeartDisease)
- ~320,000 rows

Cleaned file (2020_cleaned-~310,000 rows), python file on github-
FA_data_cleaning_2020.ipynb:

- Renamed columns to match identical columns in heart_2022 with slightly different names(ex: GenHealth:GeneralHealth; PhysicaHealth:PhysicalHealthDays; DiffWalking:DifficultyWalking,etc.)
- Renamed other columns for clarity such as SleepTime:HoursOfSleep
- Removed rows with ambiguous data: Diabetic (Yes, during pregnancy, No, borderline diabetes) to yield binary column
- Created dummy variables ideal for binary categories
- Mapped ordinal variables from least to greatest starting from 0 (GeneralHealth and AgeCategory)
- Further analysis with mapped and dummy values
- [2020 Data Cleaning Code](#)

original file (heart_2022_no_nans.csv):

- 39 feature columns (heart disease indicators)
- 1 target column ("HadHeartAttack")
- ~246,000 rows

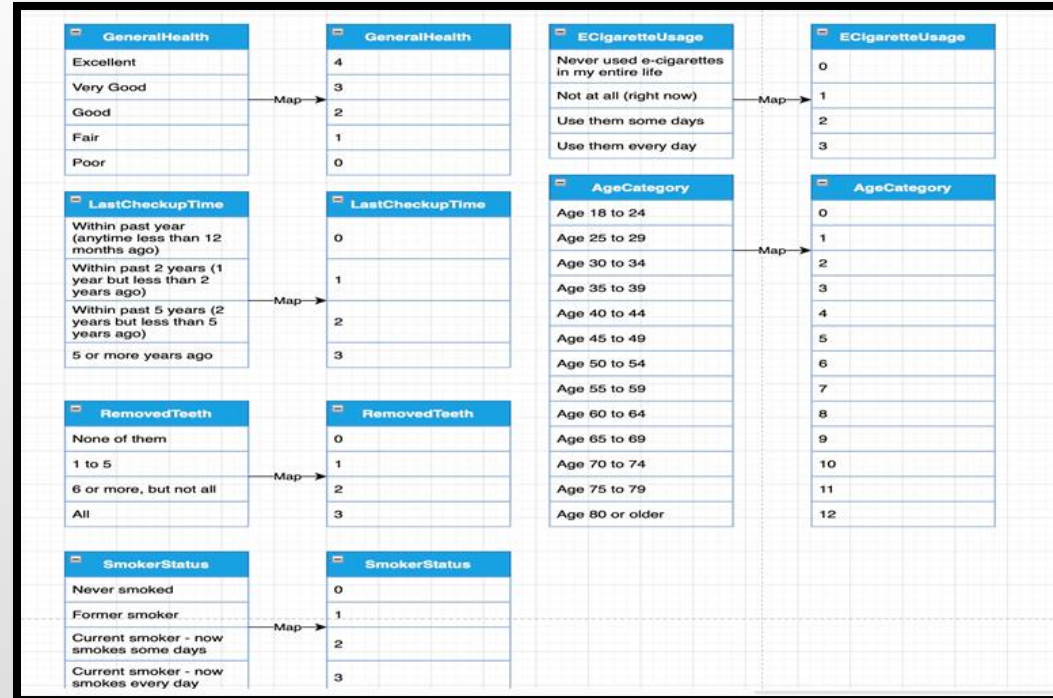
Cleaned file (cleaned_df), python file on github- data_cleaning_2022.ipynb:

- Removed columns ("State", "HadDiabetes", etc)
- Renamed other columns for clarity such as "SleepTime": "HoursOfSleep"
- Removed NaN
- Created dummy variables ideal for binary categories ("Sex", "PhysicalActivities", "HadHeartAttack", "HadAngina", etc)
- Mapped ordinal variables from least to greatest starting from 0 ("GeneralHealth", "AgeCategory", "LastCheckupTime", "RemovedTeeth", "SmokerStatus", and "ECigaretteUsage")
- [2022 Data Cleaning Code](#)

Analysis On Python Using Mapped/Dummy Variables Using Linear Regression Model

Heart Disease Indicators	r-value
0 GeneralHealth	-0.24471116023984052
1 AgeCategory	0.23425262301744743
2 BMI	0.052666342464889424
3 PhysicalHealthDays	0.17167977114662242
4 MentalHealthDays	0.028234899261603143
5 HoursOfSleep	0.009221458092622849
6 Smoking	0.10811849761589694

Heart Disease Indicators	r-value
0 Diabetic	0.18696346633644118
1 PhysicalActivities	-0.10089330845766817
2 Asthma	0.04108739544718814
3 SkinCancer	0.09320241203489989
4 AlcoholDrinking	-0.032941641681676424
5 Stroke	0.19801184782142203
6 DifficultyWalking	0.20292116027668886
7 KidneyDisease	0.14542892385601908



Heart Disease Indicators	r-value
1 GeneralHealth	-0.244711
2 AgeCategory	0.234253
3 DifficultyWalking	0.202921
4 Stroke	0.198012
5 Diabetic	0.186963
PhysicalHealthDays	0.171680
KidneyDisease	0.145429
Smoking	0.108118
PhysicalActivities	-0.100893
SkinCancer	0.093202
BMI	0.052666
Asthma	0.041087
AlcoholDrinking	-0.032942
MentalHealthDays	0.028235
HoursOfSleep	0.009221

Spark Ingestion, Data Understanding, Exporting Model

- Reading our cleaned data file from GitHub
- Data querying to get a better understanding of the data
- Exporting our data as a csv for modeling

```
import pandas as pd

heartdata2020 = 'https://raw.githubusercontent.com/oliverkiswa/Final-Project-Team-1/main/Resources/2020_cleaned.csv'

pd_df = pd.read_csv(heartdata2020)
heart2020_df = spark.createDataFrame(pd_df)

heart2020_df.limit(5).show()
```

HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealthDays	MentalHealthDays	DifficultyWalking	Sex	AgeCategory
No	16.6	Yes	No	No	3.0	30.0	No	Female	55-59
No	20.34	No	No	Yes	0.0	0.0	No	Female	80 or older
No	26.58	Yes	No	No	20.0	30.0	No	Male	65-69
No	24.21	No	No	No	0.0	0.0	No	Female	75-79
No	23.71	No	No	No	28.0	0.0	Yes	Female	40-44

```
a2020q1 = """
SELECT
  HeartDisease,
  COUNT(*) AS TOTAL,
  ROUND(COUNT(CASE WHEN Smoking = 'Yes' THEN Smoking END) / TOTAL * 100,2) AS PERCENT_SMOKING,
  ROUND(COUNT(CASE WHEN AlcoholDrinking = 'Yes' THEN AlcoholDrinking END) / TOTAL * 100,2) AS PERCENT_DRINKERS,
  ROUND(COUNT(CASE WHEN Stroke = 'Yes' THEN Stroke END) / TOTAL * 100,2) AS PERCENT_STROKE,
  ROUND(COUNT(CASE WHEN Diabetic = 'Yes' THEN Diabetic END) / TOTAL * 100,2) AS PERCENT_DIABETIC,
  ROUND(COUNT(CASE WHEN Asthma = 'Yes' THEN Asthma END) / TOTAL * 100,2) AS PERCENT_ASTHMA,
  ROUND(COUNT(CASE WHEN KidneyDisease = 'Yes' THEN KidneyDisease END) / TOTAL * 100,2) AS PERCENT_KIDNEY_DISEASE,
  ROUND(COUNT(CASE WHEN SkinCancer = 'Yes' THEN SkinCancer END) / TOTAL * 100,2) AS PERCENT_SKIN_CANCER
FROM heart20
Group by HeartDisease
ORDER BY HeartDisease DESC
"""

spark.sql(a2020q1).show()
```

HeartDisease	TOTAL	PERCENT_SMOKING	PERCENT_DRINKERS	PERCENT_STROKE	PERCENT_DIABETIC	PERCENT_ASTHMA	PERCENT_KIDNEY_DISEASE
Yes	26584	58.66	4.16	16.07	33.69	17.89	12.65
No	286430	39.58	7.09	2.6	11.12	12.9	2.84

```
[11] heart2020_df.toPandas().to_csv('2020_cleaned.csv')
```

Initial Model Creation, Running Into Issues With Precision and Prediction

- First model run with 2022 data
- “Had HeartAttack” as our target
- Removed columns that were not binary features, for better optimization
- Classification reports presented a poor precision percentage for predicting Yes for Heart Attack

```
new_df = df.drop(['AgeCategory', 'SleepHours', 'HeightInMeters', 'WeightInKilograms', 'BMI', 'H  
| | | | | 'CovidPos_No', 'PneumoVaxEver_Yes', 'HadAngina_Yes', 'HadAngina_No'], axis=1)
```

```
y = df['HadHeartAttack_Yes']  
X = df.drop(columns=['HadHeartAttack_Yes', 'HadHeartAttack_No'])  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42, stratify=y)
```

```
# Create and save the testing classification report  
testing_report = classification_report(y_test, testing_predictions)  
  
# Print the testing classification report  
print(testing_report)  
✓ 0.0s
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	37015
1	0.59	0.26	0.36	2741
accuracy			0.94	39756
macro avg	0.77	0.62	0.66	39756
weighted avg	0.92	0.94	0.92	39756



Model Problems, Reasoning For Model Changes

Initial Model Issues

- Strong when predicting who doesn't have heart disease
- Weak when predicting who does!
- Opposite of what you want

Why?

- Data was very skewed (90% didn't have heart disease)
- Model didn't have enough data to identify patients with heart disease



Understanding Issues With Health Care Data

This Is A Common Problem With Heart Disease

- We observed this issue 2020, 2022, other data sources
- Makes sense: most people don't actively have heart disease
- Heart disease seems binary but is likely not

Still A Problem!

- Clearly, our model needs work
- Not good enough to only identify who's healthy, we need to know who's sick
- We iterated through multiple common tactics to refine the model

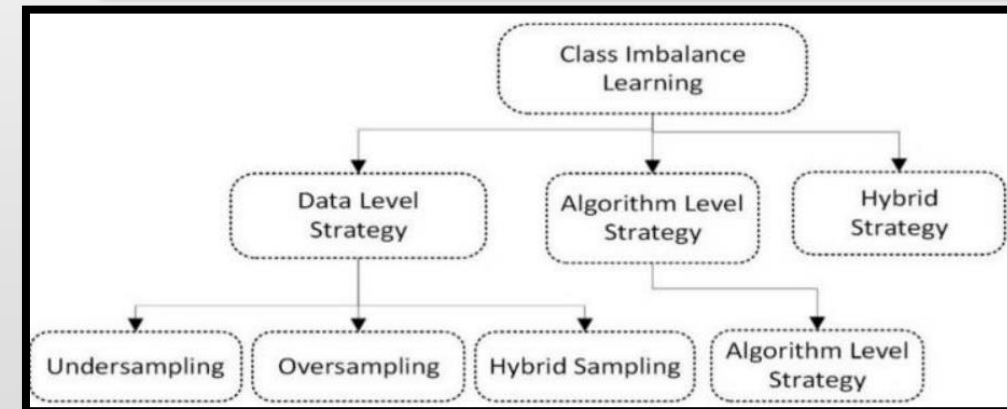
[Healthcare \(Basel\)](#), 2022 Jul; 10(7): 1293. PMID: PMC9322725
Published online 2022 Jul 13. doi: [10.3390/healthcare10071293](#) PMID: [35885819](#)

Addressing Binary Classification over Class Imbalanced Clinical Datasets Using Computationally Intelligent Techniques

[Vinod Kumar](#),¹ [Gotam Singh Lalotra](#),² [Ponnusamy Sasikala](#),³ [Dharmendra Singh Rajput](#),^{4,*} [Rajesh Kaluri](#),⁴ [Kuruva Lakshmana](#),⁴ [Mohammad Shorfuzzaman](#),⁵ [Abdulmajeed Alsufyani](#),⁵ and [Mueen Uddin](#),^{6,*}

Andrea Tittarelli, Academic Editor

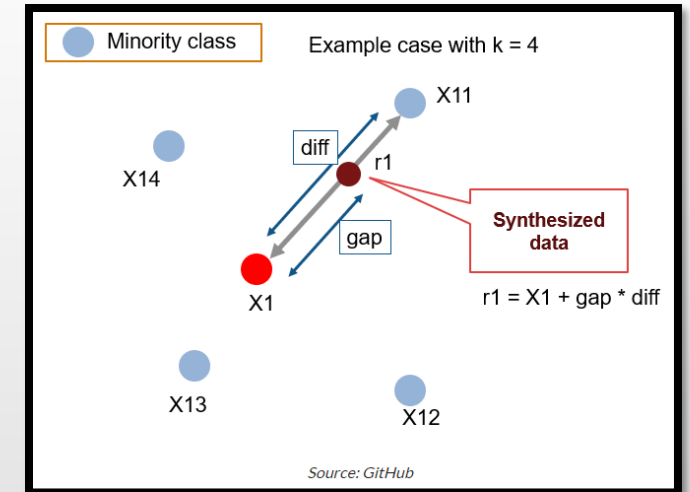
[Author information](#) [Article notes](#) [Copyright and License information](#) [PMC Disclaimer](#)



Model Optimization – SMOTE

(Synthetic Minority Oversampling Technique)

- There's An Imbalance Ratio Of No Responses For Heart Disease To Yes Responses. Causes The Accuracy Metric To Be Biased And Not Preferable.
- SMOTE Alters The Training Set By Increasing The Number Of Yes Data Points To Match The Volume Of No Data Points.
- Creates Synthetic Samples By Taking A Random Instance Of The Minority Class, Finding Its K-nearest Neighbors And Placing An New Instance At A Random Distance Between.



```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=78)

smt = SMOTE()

counter = Counter(y_train)
print(counter)

Counter({0.0: 219270, 1.0: 20576})

X_train_sm, y_train_sm = smt.fit_resample(X_train, y_train)

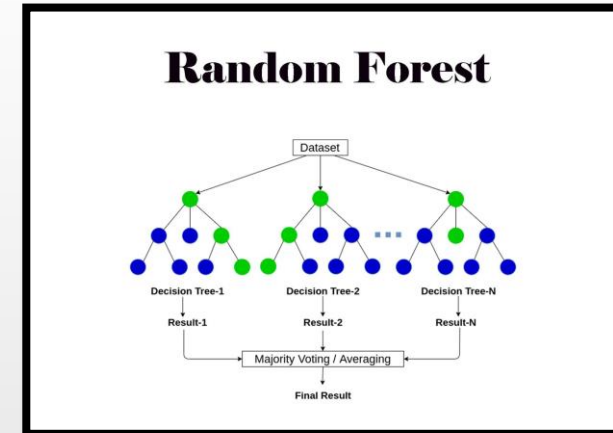
counter = Counter(y_train_sm)
print(counter)

Counter({0.0: 219270, 1.0: 219270})
```


Model Optimization – Random Forest

- Random Forests is an ensemble learning method.

Instead of one complex decision tree, it samples the data and constructs a multitude of simple decision trees.



- Unlike normal decision trees, Random Forests is robust against overfitting. SMOTE runs the risk of introducing noisy instances and overfitting problems.



```
rf_model = RandomForestClassifier(n_estimators=500, random_state=1567)

rf_model = rf_model.fit(X_train_scaled, y_train_sm)
```

Model Optimization: SMOTE + Random Forest Results

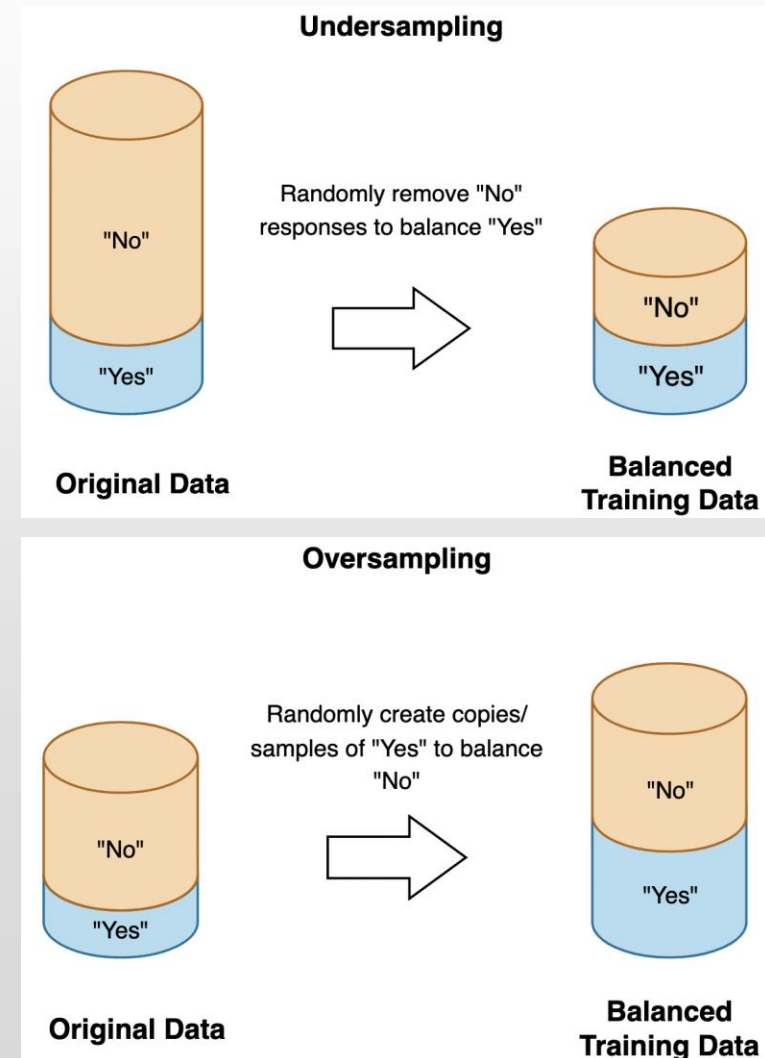
- Through SMOTE + Random Forests, the model classification report had low precision and high recall for the population with heart disease, which is preferable.
- Low precision: the model is incorrectly predicting heart disease for people who don't have it; many false positives.
- High recall: for the population that does have heart disease, the model is correctly identifying a majority of them; few false negatives.
- False positives will lead to further testing and safety precautions. A high number of false negatives may result in heart disease being untreated which might have serious consequences.

SMOTE					
Accuracy Score : 0.7500656668626249					
Classification Report					
	precision	recall	f1-score	support	
0.0	0.96	0.76	0.85	73152	
1.0	0.20	0.64	0.30	6797	
accuracy			0.75	79949	
macro avg	0.58	0.70	0.58	79949	
weighted avg	0.89	0.75	0.80	79949	

Model Optimization: Undersampling

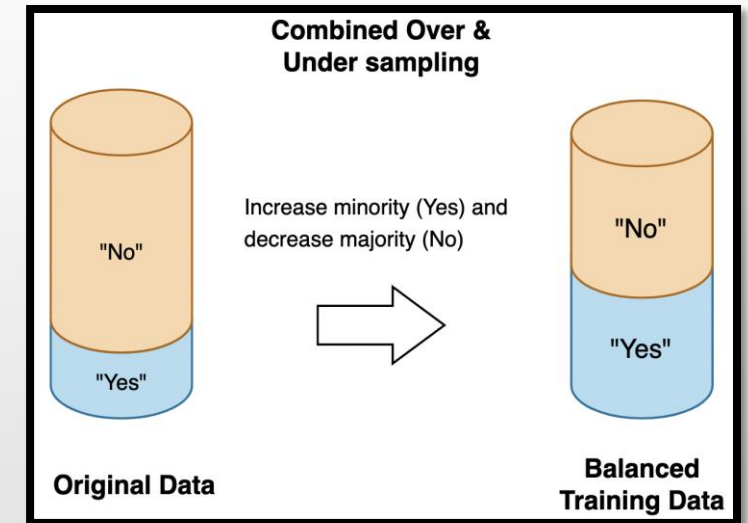
- How does it work?
 - “Opposite” of oversampling
- Undersampling model
 - Results & classification report

	precision	recall	f1-score	support
False	0.96	0.78	0.86	170404
True	0.21	0.62	0.32	15869
accuracy			0.77	186273
macro avg	0.58	0.70	0.59	186273
weighted avg	0.89	0.77	0.82	186273



Model Optimization: Combining Over & Undersampling






- Goal: “Balance” Bias Created By Either Method
 - Oversampling might add non-useful data
 - Undersampling might remove useful data
- Results & Classification Report
 - Why didn't this help much?
 - Traded off precision for recall
- Which Model Is “Best”?
 - Model became very sensitive



	precision	recall	f1-score	support
0	0.96	0.74	0.84	141962
1	0.19	0.67	0.30	13266
accuracy			0.73	155228
macro avg	0.58	0.70	0.57	155228
weighted avg	0.89	0.73	0.79	155228


Model Optimization: Combining Over & Undersampling

- Try different resampling methods
 - Rubber band analogy for weights
- Some feature columns capture overlapping information
 - Leads to unintentionally over/ under weighting variables




Further Model Improvements

Feature weights (Undersampling Model)



Feature weights (Over & Under sampling model)



Heart Disease Dashboard

Heart Disease & Race In America




Heart Disease Yes/No




Heart Disease

- ☐ (All)
- ☐ No
- ☒ Yes

Age Category & Heart Disease In America




Diabetes




Alcohol Drinking




Physical Activity




Ashtma




Skin Cancer




Smoking



Stroke Yes/No



Difficulty Walking



Thank You!